

# Limitations of Standard Autoencoders

- They struggle to learn meaningful latent spaces.
- Lack of control over latent representations.
- Enter **Variational Autoencoders (VAEs)**.

# Variational Autoencoders (VAEs)

- VAEs extend autoencoders to probabilistic models.
- Instead of a fixed latent ( $z$ ), VAEs learn a probability distribution ( $p(z|x)$ ).
- Outputs are sampled from this distribution.

# VAE Architecture

Input --> Encoder --> Latent Distribution --> Decoder --> Output

- **Encoder:** Produces  $(\mu)$  and  $(\sigma^2)$  for latent distribution.
- **Latent Distribution:** Sampled using reparameterization trick.
- **Decoder:** Reconstructs  $(x)$  from samples  $(z)$ .

# Latent Space in VAEs

- In VAEs, the latent space represents a distribution.
- Assumes  $z \sim \mathcal{N}(\mu, \sigma^2)$ , a Gaussian distribution.

# VAE Loss Function

The VAE loss combines two terms:

1. **Reconstruction Loss:**

$$L_{\text{recon}}(x, \hat{x}) = ||x - \hat{x}||^2$$

2. **KL Divergence (regularizer):**

$$L_{\text{KL}} = D_{\text{KL}}(q(z|x)||p(z))$$

# KL Divergence in VAEs

The KL divergence regularizes the latent space:

$$D_{\text{KL}}(q(z|x)||p(z)) = \int q(z|x) \log \left( \frac{q(z|x)}{p(z)} \right) dz$$

- Encourages  $q(z|x)$  to be close to the prior  $p(z)$ .

# Reparameterization Trick

To allow backpropagation through the sampling process:

- Replace  $z \sim \mathcal{N}(\mu, \sigma^2)$  with:

$$z = \mu + \sigma \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

# Kullback-Leibler (KL) Divergence

The formula for KL divergence between a reference probability distribution  $P$  and a second probability distribution  $Q$  is:

$$D_{KL}(P||Q) = \sum_{x \in \mathcal{X}} P(x) \log \left( \frac{P(x)}{Q(x)} \right)$$

- **Interpretation:** The expected excess surprise from using  $Q$  as a model instead of  $P$ .



# Reconstruction Loss


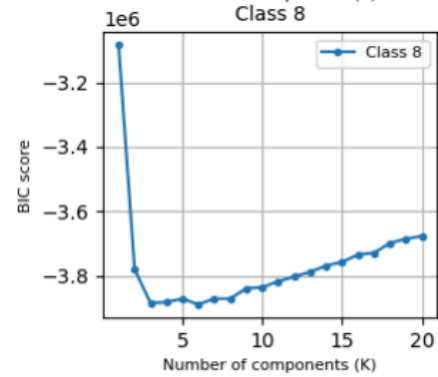


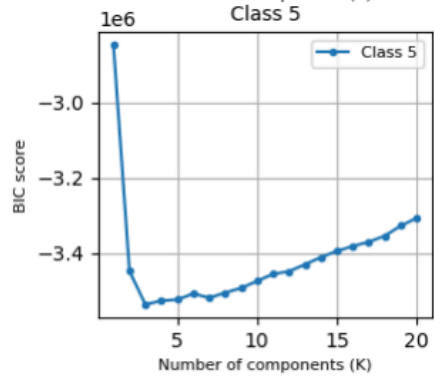

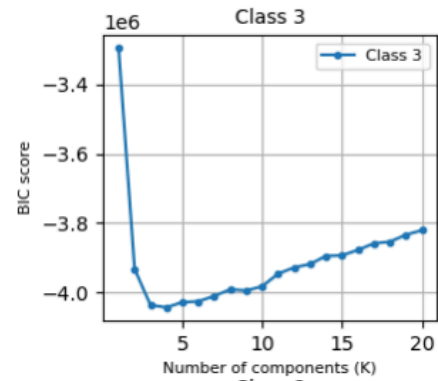

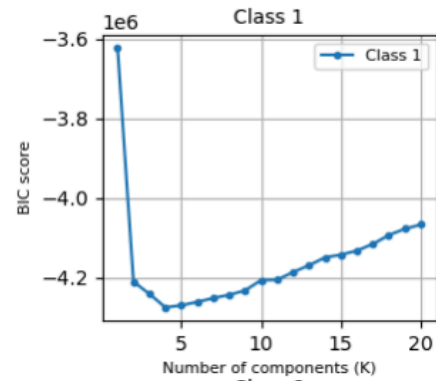

The code to compute the reconstruction loss:

```
from torch.nn import functional as F

def reconstruction_loss(x, x_hat):
    return F.binary_cross_entropy(x_hat, x, reduction="sum")
```

- Inputs: `x` and `x_hat` have dimensions  $(N, 1, H, W)$ .
- **Explanation:** This computes the negative log-likelihood of the Bernoulli distribution.

# BIC Curve



$$\frac{1}{2}$$